

Test plan - app sync in a single light loading cluster

Test environment:

- Pet cluster created by cluster orchestrator
- Deploy new argocd
- In-cluster self-host argocd: <https://github.com/J Julian-Chu/argocd-memory-optimization-test/tree/main/argocd>
- Guestbook app * 41 (synced): <https://github.com/J Julian-Chu/argocd-memory-optimization-test/tree/main/guestbook>
- No CPU and Memory request and limitation
- Keep 2 hours no configuration changes to collect metrics, sync the 41 guestbook apps per hour

Test case	Controller status processors	Controller operation processors	GOMEMLIMIT	Resource inclusion	Time(GMT+2)	Config changes
1	20(default)	10(default)	No	false	2025/09/16 1900-21:00 Sync1: 19:30 Sync2: 20:30	https://github.com/J Julian-Chu/argocd-memory-optimization-test/commit/0c9a060c213c62dff77fd1edd9290a02d28e9ec
2	2	10	No	false	2025/09/16 22:00-24:00 Sync1: 22:30 Sync2: 23:30	https://github.com/J Julian-Chu/argocd-memory-optimization-test/commit/059872a650bef0c27556c5f64a815457d9e9d87c
3	2	1	No	false	2025/09/17 11:00-13:00 Sync1: 11:30	https://github.com/J Julian-Chu/argocd-memory-optimization-test/commit/db065f01691c849c481f

					Sync2: 12:30	395fa2ae24042784122a
4	20	10	500MiB	false	2025/09/17 14:00-16:00 Sync1: 14:30 Sync2: 15:30	https://github.com/Julian-Chu/argocd-memory-optimization-test/commit/efd919a9a8e75a39dbe8506991b9c219ad073d97
5	20	10	No	yes	2025/09/17 18:00-20:00 Sync1: 18:30 Sync2: 19:30	https://github.com/Julian-Chu/argocd-memory-optimization-test/commit/3daaf0e650c08abd08099768ee763e0a00ec7ee3

Test case 1: (default setting, baseline)



CPU:

- Normal: < 0.1
- Spike after sync: 0.7

Memory:

- Normal: 384 - 765 MiB
- Spike after sync: 1003 MiB

Test case 2: (reduce status processors from 20 to 2)



CPU:

- Normal: < 0.1
- Spike after sync: 0.7

Memory:

- Normal: 403 - 740 MiB
- Spike after sync: 1000 MiB

Summary:

No obvious changes

Test case 3: (reduce operation processors from 10 to 1)



CPU:

- Normal: < 0.1
- Spike after sync: 0.667

Memory:

- Normal: 396 - 714 MiB
- Spike after sync: 998 Mib

Summary:

No obvious changes

Test case 4: (GOMEMLIMIT: 500MiB)



CPU:

- Normal: < 0.12
- Spike after sync: 2.24

Memory:

- Normal: 423 - 516 MiB
- Spike after sync: 520 Mib

Summary:

- The memory usage remain stable near the GOMEMLIMIT value
- CPU utilization spikes to **three times its normal baseline** during peak load.
- The CPU spike is expected because the garbage collector can be triggered more aggressively

Test case 5: (add resource.inclusion)



CPU:

- Normal: < 0.1
- Spike after sync: 0.645

Memory:

- Normal: 391 - 706 MiB
- Spike after sync: 957 MiB



Argocd managed api resources:

- 672 → 241
- Objects : 4200 → 3000 (still high)

Summary:

- No obvious changes in the CPU and memory, but huge reduction in the resources.
- Reducing the number of resources included should lower memory usage. However, we're not seeing that here, possibly because the guestbook only has few resources.
- 3000 is still a high number, any opportunity to reduce it ?

Summary:

After running tests in the light loading cluster, our tests indicate that `GOMEMLIMIT` effectively reduces memory usage and spikes in lightly loaded clusters, though it causes a temporary CPU spike. Given that CPU is a burstable resource and the ArgoCD application controller is not latency-sensitive, these short-term spikes are acceptable.

We propose setting `GOMEMLIMIT` to 90% of the memory limit to mitigate memory spikes, while closely monitoring CPU and sync performance.

More information:

- **Single light loading cluster:** Observed a peak memory usage of 1 GB.
- **Impact of `resource.inclusion`:** While `resource.inclusion` reduced the number of ArgoCD-managed API resources, no corresponding decrease in memory usage was observed. This limited impact might be attributed to the test's scope, as the guestbook application uses only a single pod and may not accurately represent real-world scenarios.
- **Status and operation processor:** These components had no noticeable impact, likely because the applications were consistently synchronized, eliminating the need to apply new resources.